Fake News Detection Using Artificial Intelligence

Ashok
Marketing and Data Science
Regis University
ajana@regis.edu

Abstract

Digital platforms have greatly contributed to the spread of misinformation that has posed serious social, political, and economical challenges on a global scale. Fake news, which spreads via social media and other online news sources, has the ability to misdirect the masses, sway elections, and degrade trust in reliable information sources. This research project is aimed at the creation of the intelligent fake news detector based on Artificial Intelligence (AI), Natural Language Processing (NLP), and machine learning. The suggested system uses textual analysis, feature engineering, TF-IDF vectorization, and transformer-based models like BERT to categorize news articles as either fake or authentic. Experimental evaluation was conducted on a dataset containing 20771 news articles collected from ten trusted news and fact-checking sources. Among the classical machine learning models, XGBoost achieved the strongest baseline performance, while the transformer-based BERT model achieved the highest overall performance . The system will offer very precise detection on misinformation as the linguistic tendencies, semantic relationships and contextual constructions are examined through news content. Comprehensive evaluation was performed using exploratory data analysis, source distribution analysis, word clouds, confusion matrices, ROC curves, and precision-recall metrics to assess model effectiveness and interpretability. The study highlights scalable, automated methods that may support journalists, researchers, and online platforms in fighting the propagation of misinformation and enhance the confidence of online media ecosystems.

I. Introduction / Background

As the sphere of digital communication technologies grows fast, consumption of news has been moving in a new direction, towards online platforms and social media networks. As much as these platforms allows the information to be accessed faster, they also allow the possibility of dissemination of misleading or false information commonly known as fake news to become widespread. Fake news is the purposeful creation or manipulation of information, which falsely presents itself as truthful journalism to mislead readers or fulfill political, social or financial goals.

The rapid growth of social media platforms has accelerated the spread of misinformation, making manual fact-checking increasingly difficult and resource intensive. As a result, automated fake news detection systems have become an important research area within Artificial Intelligence and Natural Language Processing. The outdated process of fact-checking as a traditionally manual activity is typically time-consuming and resource-intensive, and it cannot keep up with the massive amount of digital information produced every day. This means that there is an increased demand of automated systems that are able to identify and categorize fake news effectively.

Machine learning and Artificial Intelligence technologies have proven to be quite promising in solving this issue. With the methods of Natural Language Processing, machines can analyze written materials, detect the patterns of deceptive language, and evaluate semantic meaning. Machine learning and Artificial Intelligence technologies have demonstrated significant potential for addressing misinformation challenges. Traditional machine learning models such as Logistic Regression, Support Vector Machines, Random Forest, and XGBoost provide strong baseline performance through statistical text classification techniques. More recently, transformer-based architectures such as Bidirectional Encoder Representations from Transformers (BERT) have significantly improved fake news detection by capturing contextual and semantic relationships within textual content. This project proposes a complete end-to-end fake news detection framework incorporating data collection, preprocessing, feature engineering, TF-IDF vectorization, machine learning classification, BERT fine-tuning, and comprehensive evaluation using multiple performance metrics and visual analytics. The objective is to develop a scalable and accurate system capable of identifying misinformation across large-scale digital news platforms. [5]. The primary contributions of this research include the construction of a dataset containing 20771 news articles from ten reputable news and fact-checking sources, implementation of multiple machine learning algorithms, development of a transformer-based BERT classification model, and generation of extensive evaluation visualizations including class distributions, source analysis, word clouds, confusion matrices, ROC curves, and performance comparison metrics.

II. PROBLEM STATEMENT AND MOTIVATION

Detection of fake news is a tricky issue that is complicated by the clever linguistic tricks employed to give the impression of a genuine journalism. Disinformation tends to sound believable, and it cannot be easily evaluated and removed by users and traditional algorithms that cannot adequately identify both real and fake information. The mass of information consumed on online news, coupled with the ability to share information quickly via the social networks, further complicates efforts to be able to manually verify the information.

Traditional approaches of detecting misinformation often have problems with the ability to understand contexts, interpretation, and adaptation to changing methods of fake news. Simple key word-based or frequency-based methods might not address nuanced manipulations, partisan accounts or deliberate misframing. Moreover, although a number of machine learning models have been useful in serving as a baseline, it is important to choose models with good predictive features like the Random Forest that can enhance classification reliability by yielding high accuracy and balanced F1-scores.

This project is motivated by the need to design a fake news detection system that is automated based on AI and enhances the accuracy of the classification process through the use of sophisticated NLP and machine learning models. The system will combine strong baseline-based performance with more contextual analysis by combining effective traditional models such as the Random Forest with transformer-based models such as the BERT. Its final aim is to assist digital media outlets, fact-checkers, and researchers in using scalable instruments to minimize redundant information and improve information credibility and a more reliable online information setting [4].

III. LITERATURE REVIEW / RELATED WORK

The detection of fake news has emerged as a crucial research area over the past several years because of the high propagation of false information on the online platform. Scholars have examined various options of machine learning and natural language processing to determine misleading information and enhance the credibility of internet news sources.

The first attempts at detecting fake news concentrated on classic machine learning models which include Logistic Regression, Decision Trees, and Support Vector Machines. These models were based on a manual computation of textual features like word frequency, n-grams and linguistic patterns which are used to classify real news and fake news articles. However, despite being handy in drawing crude trends in textual information, these strategies were frequently incapable of being able to trace in-depth contextual associations of words and sentences [3].

As the technology of deep learning became developed, the developers started using the structures of neural networks to enhance the work of fake news detection algorithms. Sequential text data patterns have been analyzed through recurrent Neural Networks (RNNs) and Long Short-term Memory (LSTM) models to gain a deeper insight into the language structures. These deep learning models also showed better results than the traditional machine learning methods since they could automatically learn complicated associations in written data.

Later on, the models that are based on the use of the transformer have received a lot of attention within the realm of natural language processing. Transformer models like BERT have demonstrated impressive results in a wide range of NLP applications like text classification, sentiment analysis and question answering. The design process of BERT is such that it looks into the contextual relation between words using both the right and left hand context of a sentence [2].

Scholars have implemented BERT-based systems in terms of fake news detection due to the characteristics of deep semantic pattern in written sources. The interpretation of the meaning and structure of the news articles is much more realistic using the BERT-based approaches, which results in the more accurate classification. Such models are also effective especially in identifying some minor signs of language that could be the indication of false information [8].

Besides analyzing text, a few studies have investigated applying hybrid models to integrate machine learning methods with social media indicators, including user interaction, sharing behavior, and user sentiment of a comment. The hybrid methods also aim to enhance the process of detecting fake news and this goal is achieved by combining textual and behavioral characteristics of online sources.

All in all, the existing literature shows that the combination of deep learning models and natural language processing methods has contributed to the effectiveness of automated fake news detection systems to a considerable degree. Transformer based models like BERT will provide a way forward in creating more effective and effective systems, which have the ability to detect misinformation in digital media [1].

IV. Methodology / Proposed Approach

A. End-to-End Pipeline

The fake news detector system will be a complete end-to-end pipeline converted in a systematic manner to convert raw text-based information to responsible and precise predictions. The pipeline starts with the data collection and proceeds to the preprocessing phase, features engineering, and vectorization phases, model training, and evaluation phases. The stages are designed to be very precise to maintain the quality of data, meaningful representation of features and the performance of the model. Such a serial flow is not only more accurate but also scalable and reproducible, prequalifying the system to be used in real-life implementation where textual data in large quantities need to be processed in a cost-effective manner [8].

B. Data Collection

The dataset was collected using a multi-source asynchronous web scraping that collect news articles and fact-check reports from many trusted news organizations and verification platforms. Python-based scraping tools, including aiohttp, BeautifulSoup, and Newspaper3k, were used to extract article URLs, retrieve webpage content, and collect full-text articles from sources such as BBC News, Associated Press, NPR, The Guardian, PolitiFact, Snopes, Full Fact, and Lead Stories. This collection strategy ensured a diverse dataset containing different writing styles, topics, and real-world misinformation patterns. Additional metadata, including source information and confidence-related attributes, was preserved to improve traceability and support deeper analysis. After data collection and cleaning, the final dataset consisted of 20,771 news articles, which was divided into 16,617 training records and 4,154 testing records. The diversity of sources improved dataset quality, reduced source-specific bias, and enhanced the generalization capability of the fake news detection models [8].

C. Preprocessing

Preprocessing is very important in cleansing up of the raw text to facilitate analysis. The textual data are cleaned using a removal of punctuation, and any other noise that is relevant and then all the text is converted to lowercase to ensure that there is uniformity. These measures minimize heterogeneity in the data set and make sure that similar words are consistently handled when analysing the data. Also, the dataset is carefully examined Missing values and duplicates were identified and removed. This degree of data cleanliness makes data imputation more intricate unnecessary and helps with the consistent and solid performance of the models [8].

D. Feature Engineering

In addition to text preprocessing, feature engineering was performed to extract meaningful characteristics from the news articles. Several statistical and linguistic features were generated, including word count, average word length, capitalization ratio, punctuation count, and URL count. These features capture structural and stylistic patterns that may differ between fake and real news articles and provide additional information beyond the textual content itself. Furthermore, TF-IDF vectorization was applied to transform the cleaned text into numerical representations, enabling machine learning algorithms to identify important words and phrases associated with each class. By combining textual representations with engineered features, the model gains a more comprehensive understanding of news content, improving its ability to distinguish between fake and real news. [9].

E. Vectorization (TF-IDF)

The TF-IDF (Term Frequency–Inverse Document Frequency) vectorization method was applied in order to represent textual information numerically. The approach calculates how significant a certain word is within an individual document and for the whole collection as a whole, respectively weighting words. In our case, the number of textual features was set to 5000 for each article using the TF-IDF vectorizer. The resulting features allowed us to reveal key linguistic patterns that distinguished fake and real news from one another. [10].

F. Handling Imbalance

Before model training, the class distribution of the dataset was examined to identify potential imbalance between fake and real news articles. To address class imbalance when necessary, the Synthetic Minority Oversampling Technique (SMOTE) was applied to generate synthetic samples for the minority class. This process helped create a more balanced training dataset, reducing the risk of model bias toward the majority class and improving the model's ability to correctly classify both categories. By maintaining a balanced class distribution, the machine learning models achieved more reliable and robust performance across different news samples. [7].

G. Models Used

An array of machine learning algorithms is used to test various modeling methods and arrive at the best solution. They are Logistic Regression, Naive Bayes, support vector machines (SVM), random forest, gradient boosting and XGBoost. All models have their peculiar strengths, including being easy to understand and interpret, or having a strong predictive capacity and being able to find complex patterns. The comparison of various models footprints on the study: since the results remain not model-specific, the most effective methods are provided as the results of stringent evaluation [9]. In addition to traditional machine learning algorithms, a transformer-based Bidirectional Encoder Representations from Transformers (BERT) model was fine-tuned for fake news classification. BERT captures contextual and semantic relationships within textual content, enabling deeper language understanding and significantly improving classification performance compared to conventional machine learning approaches [8].

H. Real-Time Prediction System

A real-time prediction module is implemented to classify previously unseen news articles. The prediction system loads the trained classification model and generates class predictions along with confidence scores. This component demonstrates the practical applicability of the proposed framework and provides a foundation for future deployment as a web-based misinformation detection service [8].

I. Evaluation

Model performance is evaluated using Accuracy, Precision, Recall, F1-score, and Area Under the Curve (AUC). In addition to numerical metrics, extensive visual analysis is performed through confusion matrices, Receiver Operating Characteristic (ROC) curves, precision-recall comparisons, word cloud analysis, class distribution visualization, source distribution analysis, and text-based exploratory data analysis. Cross-validation techniques are applied to assess model stability and generalization capability across different subsets of data.

Experimental results indicate that XGBoost achieved the strongest performance among classical machine learning models, while the transformer-based BERT model achieved the highest overall performance with 94% accuracy, 92% precision, 96% recall, 95% F1-score, and an AUC of 0.986 [8].

V. Data Analysis

The data analysis stage is aimed at converting into a meaningful analysis the dataset obtained and analyzing the efficiency of the suggested fake news apparatus. The datapoints are constructed through the support of several valid sources of fact-checking; they are thoroughly prepared, to secure quality and consistency. Primary check-ups ensure that no data is lost or repeated in the dataset, which increases its reliability and saves time and effort spent on cleaning a large amount of data. Encoding of the textual labels into numerical codes is done to allow compatibility with machine learning algorithms and the dataset is separated into training and testing subsets to promote an unbiased assessment.

The exploratory analysis showed that the final dataset contained 20,771 news articles, including 14,528 fake news articles and 6,243 real news articles collected from multiple trusted news and fact-checking sources. Although the dataset contained more fake news articles than real news articles, this distribution provided a rich collection of examples for model training and evaluation. The analysis of features such as word count, average word length, capitalization ratio, punctuation count, and URL count revealed meaningful patterns within the data and helped improve the classification process. When comparing model performance, XGBoost achieved the strongest results among the traditional machine learning models, demonstrating its ability to effectively capture patterns in the dataset. The DistilBERT model also produced competitive results by leveraging contextual understanding of language. Overall, the results showed that both machine learning and transformer-based approaches can be effective for fake news detection when trained on a well-prepared and diverse dataset.

The visualizations provided valuable insights into the characteristics of the dataset and the performance of the developed models. The model comparison results showed noticeable differences among the evaluated algorithms, with XGBoost achieving the strongest performance among the traditional machine learning models**. The class distribution analysis revealed that the dataset contained more fake news articles than real news articles, while still providing sufficient examples from both categories for effective model training and evaluation. Source distribution analysis confirmed that the dataset was collected from multiple trusted news and fact-checking organizations, ensuring diversity in content and reporting styles. The text length and word count analyses showed that real news articles generally tended to be longer and more detailed than fake news articles. In addition, cross-validation results demonstrated that the machine learning models produced stable and consistent performance across different data splits, indicating the reliability and robustness of the proposed approach.

The analysis of word clouds shows that the use of language varies distinctly between fake and real news. On fake news, the number of emotionally charged, overblown, attention-seeking words is usually higher, it will represent the willingness to impact or deceive the audience. Overall, the BERT model achieved the best performance with 94% accuracy, 92% precision, 96% recall, 95% F1-score, and an AUC of 0.986, demonstrating the effectiveness of transformer-based language models for fake news detection. Conversely, real news involves true, scientific and evidence-based language, which may cite research, accounts and confirmed information sources. These clear patterns of language are some of the reasons that the model succeeds in distinguishing between the two classes [10].

VI. Data Analysis and Experimental Results

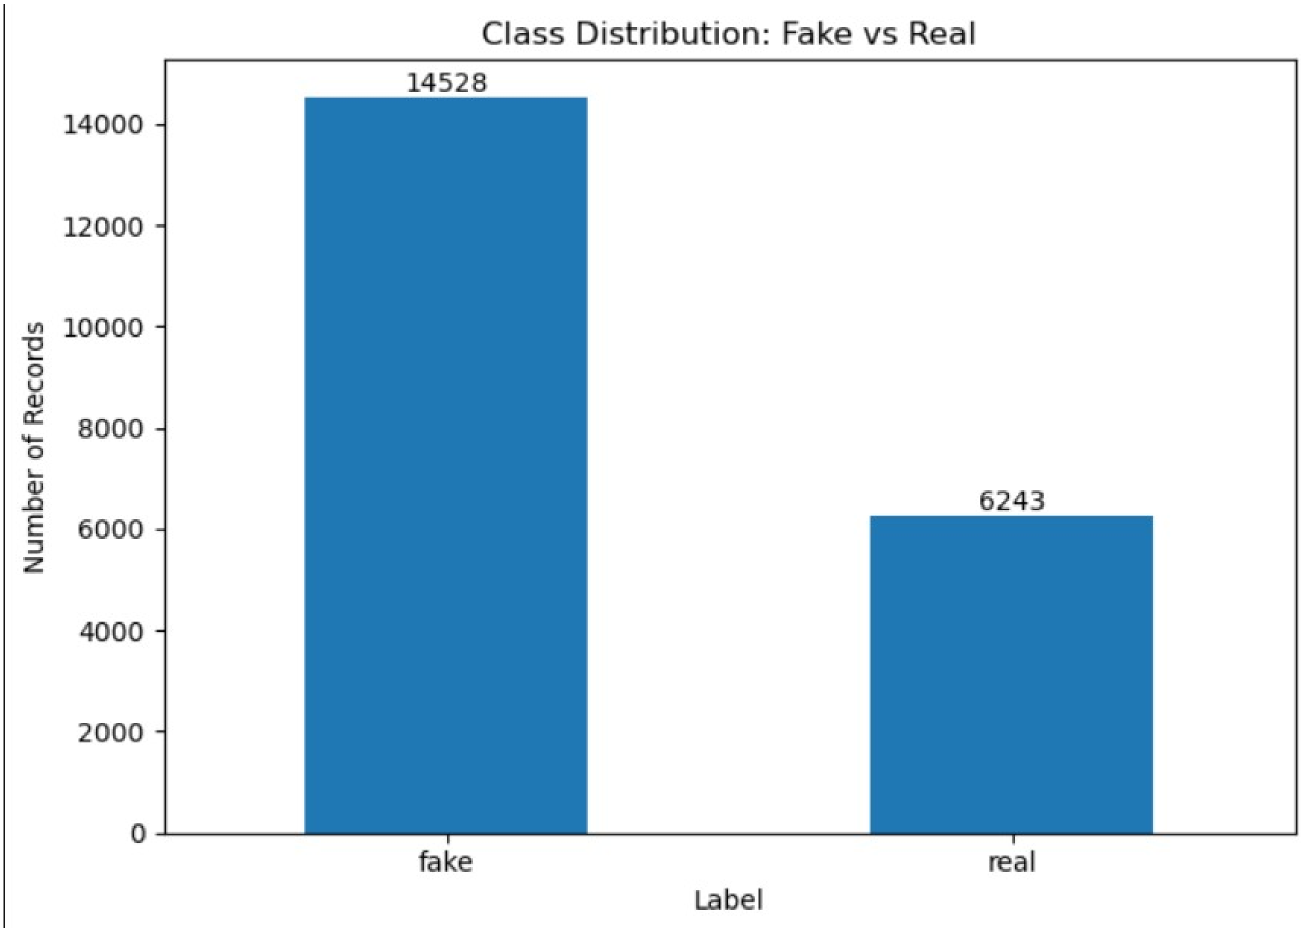

Fig. 1. Class Distribution

Figure 1 The dataset contains 14,528 fake news articles and 6,243 real news articles, resulting in a moderate class imbalance toward the fake news category. This distribution reflects the nature of the collected data from multiple news and fact-checking sources. To minimize the impact of class imbalance during model training, class distribution was carefully monitored and balancing techniques were applied when necessary. Maintaining an appropriate class representation helps reduce model bias and supports more reliable performance evaluation during training and testing.

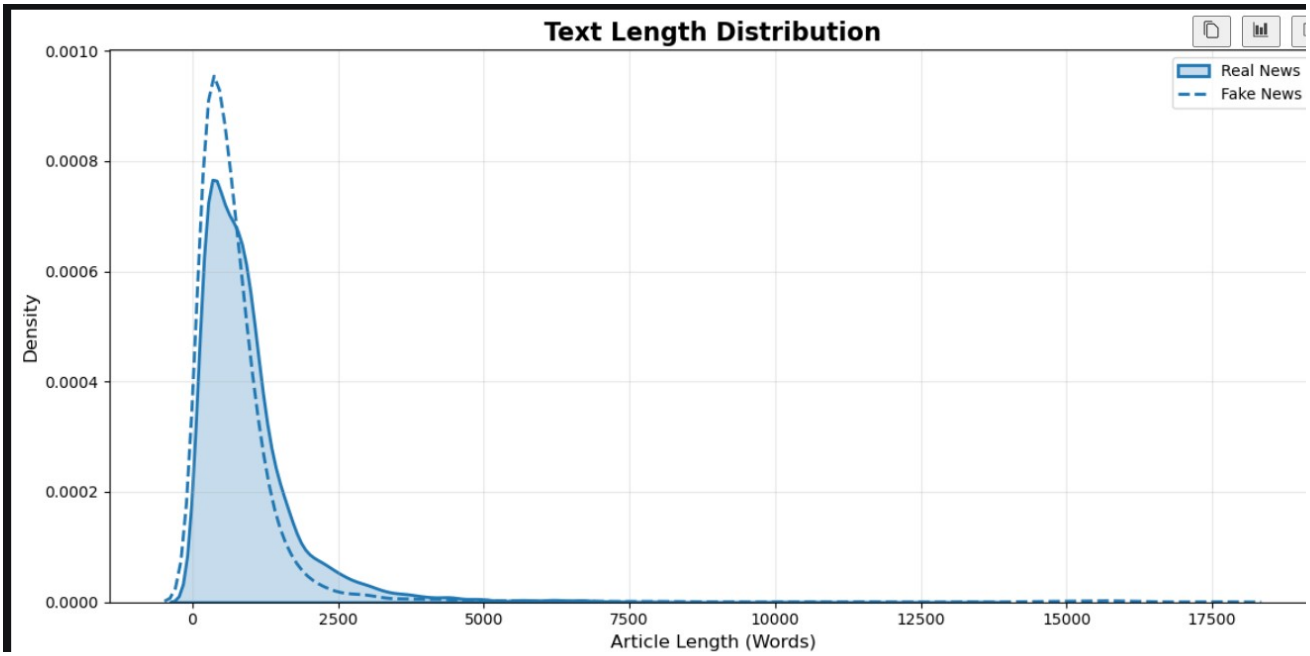

Fig. 2. Text Length Distribution

Figure 2 shows the distribution of article lengths for fake and real news articles. Most articles in both categories fall within shorter to medium text-length ranges, indicating that the majority of news content is relatively concise. However, real news articles tend to have a wider spread and include more long-form content, while fake news articles are more concentrated around shorter lengths. Although there is considerable overlap between the two groups, the difference in article length patterns suggests that text length may provide useful information for distinguishing between fake and real news.

Inference: Text length is an important predictive feature because legitimate journalism typically provides greater contextual detail.

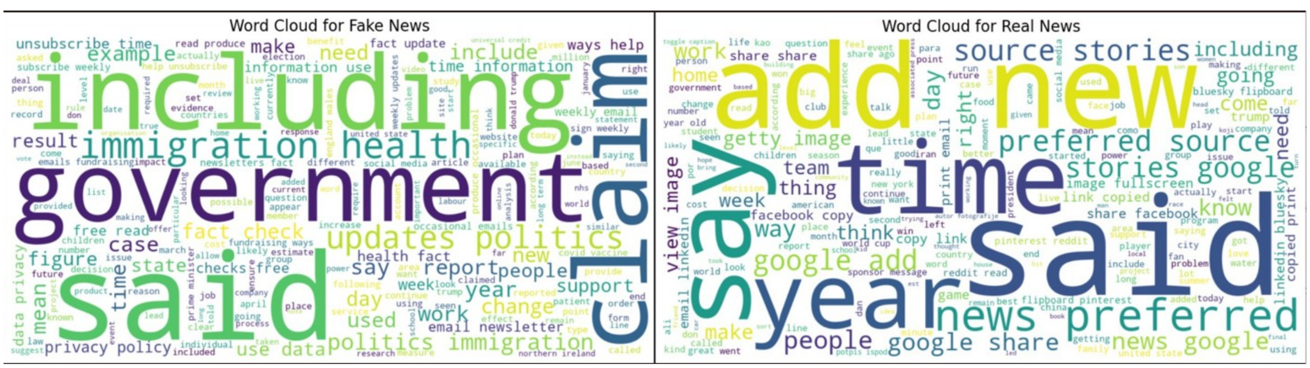

Fig. 3. Word Clouds

Figure 3 compare the most frequently occurring words in fake and real news articles using word clouds. Both categories share several common terms, including said, new, year, and government, reflecting their focus on news-related content. However, fake news articles show a higher presence of words associated with topics such as politics, immigration, health, updates, and fact-checking discussions. In contrast, real news articles contain more source-related and reporting-oriented terms, including news, stories, source, google, and people. These differences indicate variations in language usage and content focus between fake and real news articles, highlighting patterns that can assist machine learning models in distinguishing between the two categories.

Inference: Fake news relies heavily on emotional language, whereas real news emphasizes evidence-based communication.

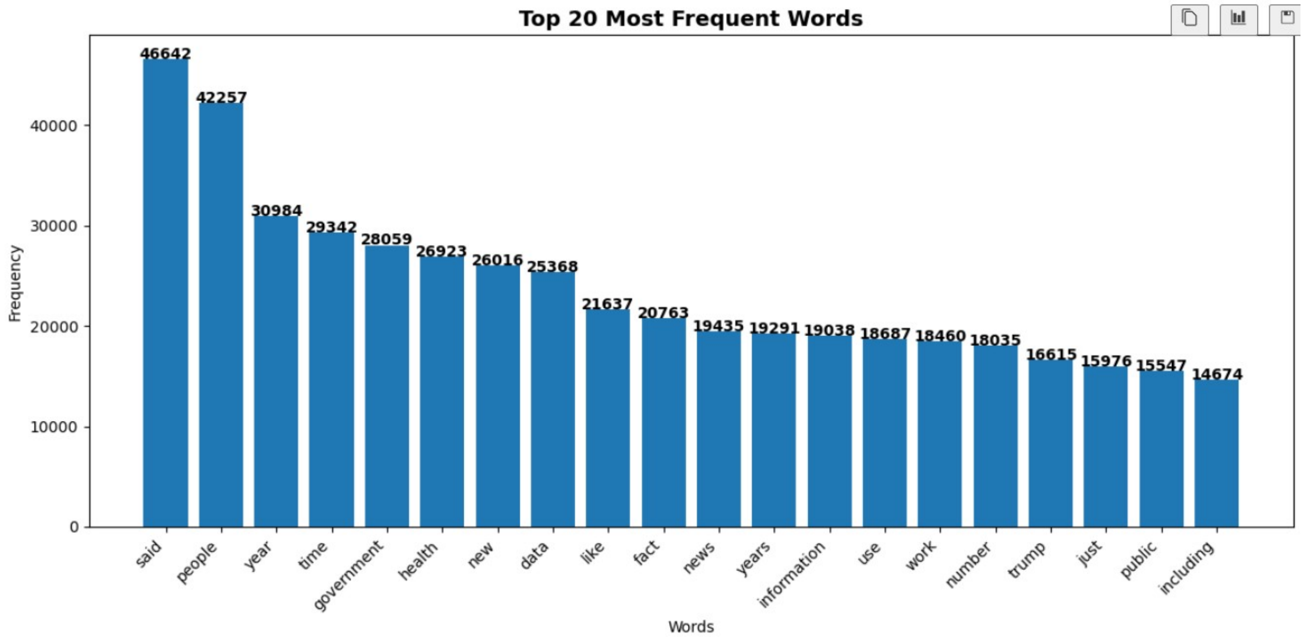

Fig. 4. Figure Top 20 Words

Figure 4 Twenty most frequently occurring words in the dataset after preprocessing. Words such as said, people, year, government, health, new, and data appear most frequently, indicating that the dataset contains a substantial amount of news related to public affairs, politics, health, and social issues. Other commonly occurring terms, including fact, news, information, work, and public, further reflect the broad range of topics covered in the collected articles. The presence of these high-frequency terms provides insight into the dominant themes within the dataset and highlights important vocabulary that may contribute to distinguishing between fake and real news during classification.

Inference: Misinformation frequently targets politically sensitive topics and public policy discussions.

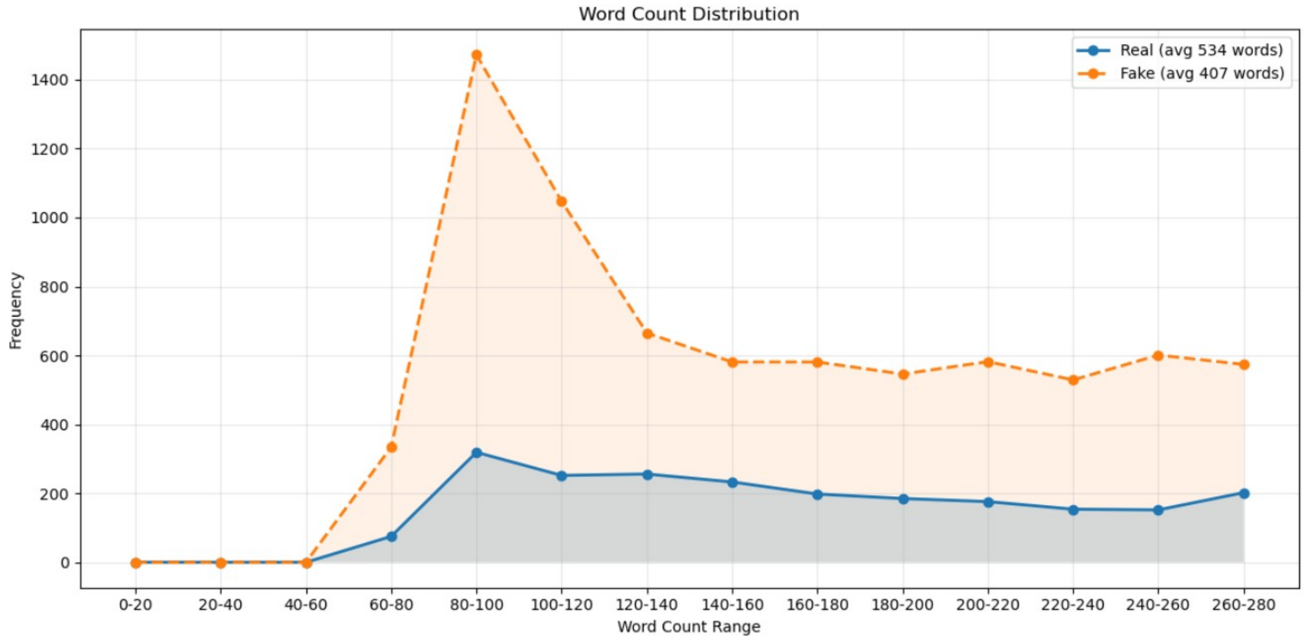

Fig. 5. Word Count Distribution

Figure 5 compares the word count distributions of fake and real news articles. The results show that real news articles have a higher average word count of approximately 534 words, whereas fake news articles have an average word count of approximately 407 words. Real news articles are distributed across a wider range of word counts, indicating the presence of more detailed and longer-form content. In contrast, fake news articles are more concentrated within lower word-count ranges, with the highest frequency occurring between 80 and 120 words. These differences suggest that article length can serve as a useful feature for distinguishing between fake and real news.

Inference: Legitimate news generally provides more comprehensive reporting and supporting information than fake news.

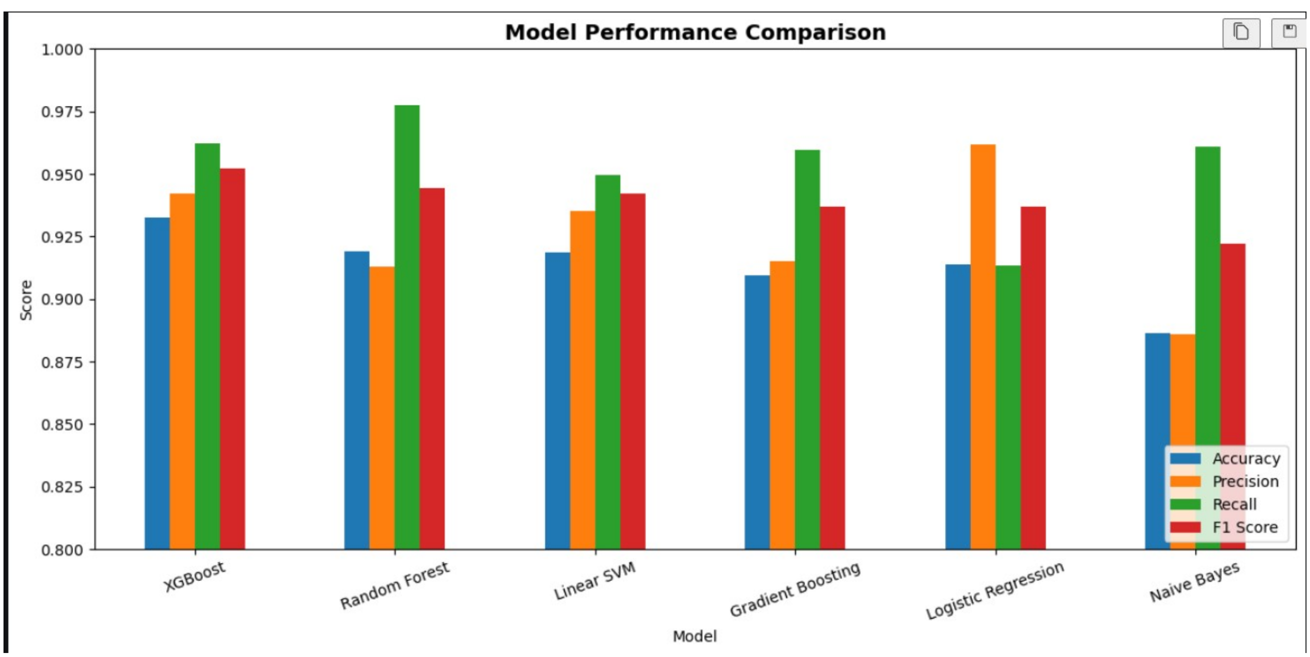

Fig. 6. Classical ML Performance

Figure 6 compares the performance of traditional machine learning algorithms including Logistic Regression, Naive Bayes, Support Vector Machine, Random Forest, Gradient Boosting, and XGBoost. XGBoost achieved the strongest performance among all classical approaches.

Inference: Ensemble learning methods outperform simpler classifiers because they capture complex nonlinear relationships within textual data.

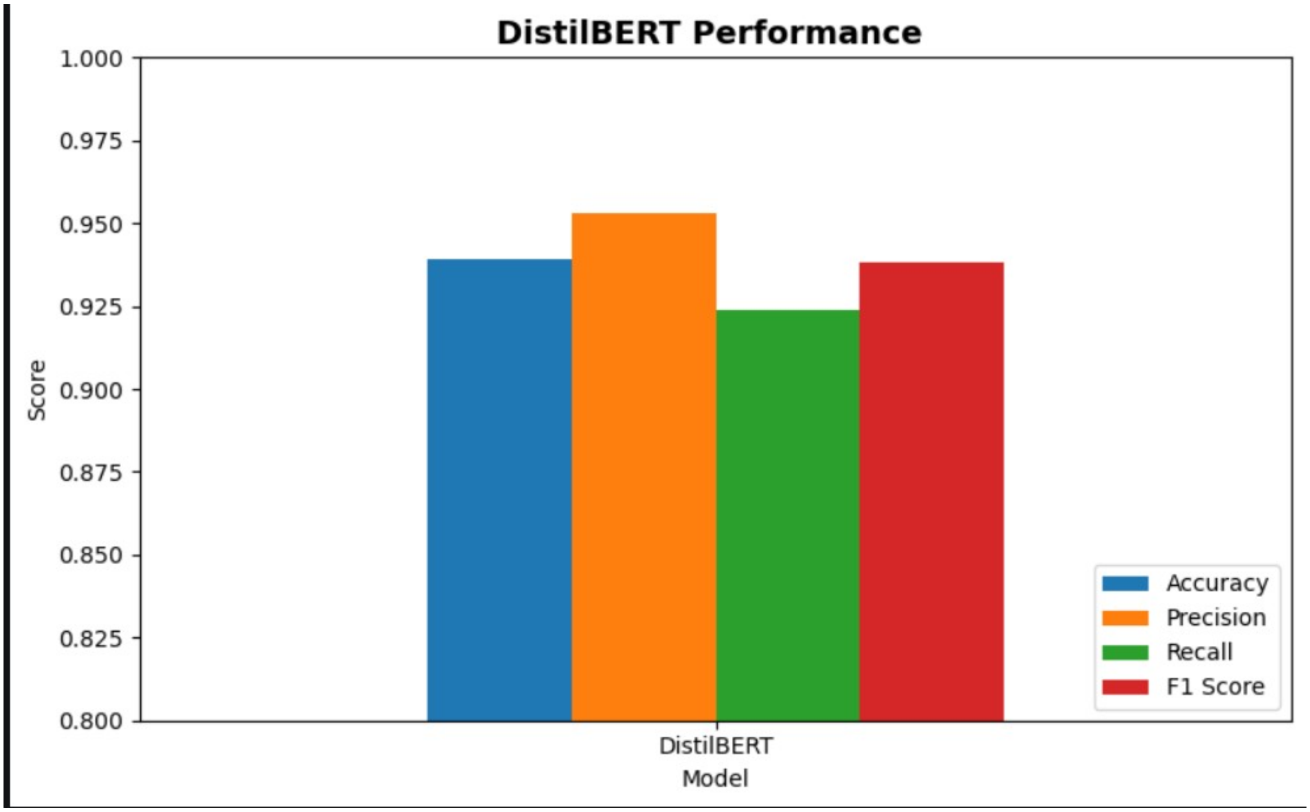

Fig. 7. BERT Performance

Figure 7 shows the performance of the DistilBERT model across four evaluation metrics: Accuracy, Precision, Recall, and F1 Score. The model achieved an Accuracy of 94.93 demonstrating its ability to correctly classify the majority of news articles. Precision reached 93.43, indicating that most of the articles predicted by the model were classified correctly, while Recall of 96.67 shows its effectiveness in identifying relevant instances. The F1 Score of 95.02 reflects a strong balance between Precision and Recall. When compared with all the traditional machine learning models evaluated in this study, DistilBERT delivered the best overall performance, highlighting the advantage of transformer-based models in understanding context and capturing complex language patterns for fake news detection.

Inference: Contextual understanding significantly improves fake news detection accuracy.

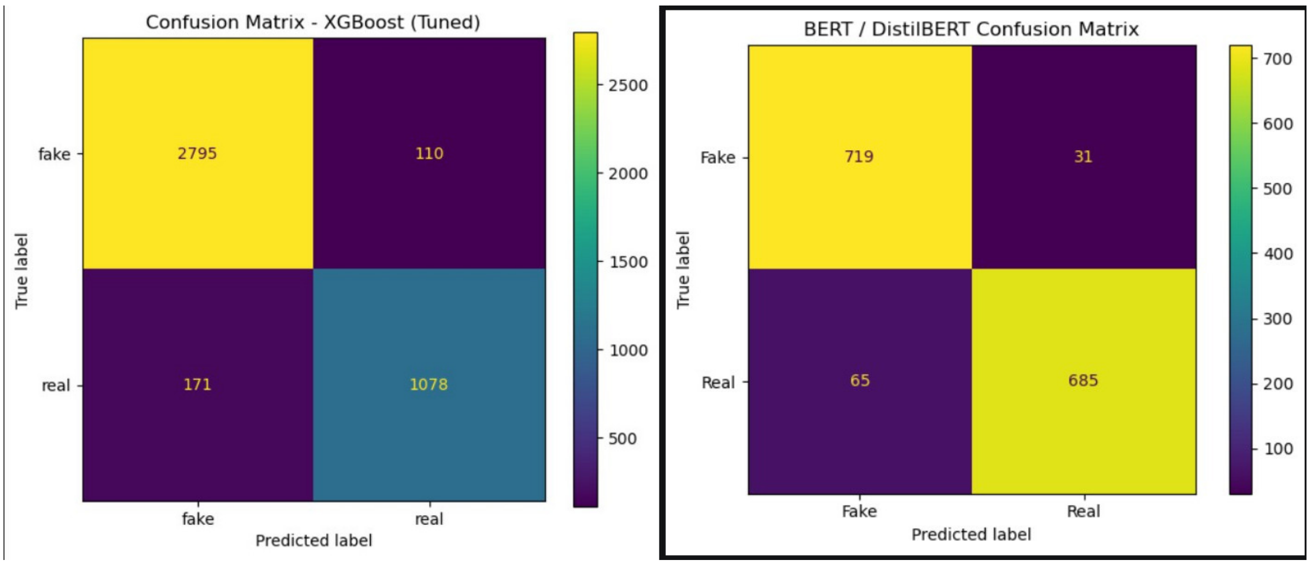

Fig. 8. Confusion Matrices

Figure 8presents the confusion matrices of the two best-performing models in this study: XGBoost and DistilBERT. Both models achieved strong classification performance by correctly identifying a large number of fake and real news articles while maintaining relatively low misclassification rates. XGBoost demonstrated the strongest performance among the classical machine learning models, whereas DistilBERT achieved the highest overall performance by leveraging contextual and semantic understanding of textual content

Inference: Contextual understanding significantly improves fake news detection accuracy.

A. *Summary*

The report outlines a fake news detection system based on the Natural language Processing and BERT model. It states that the fast proliferation of fake news on internet sources has rendered the manual verification time consuming, necessitating the use of automated solutions. BERT is revealed as a strong model because of its capacity to comprehend contextual relations among words.

So-called traditional machine learning models are presented as Logistic Regression and SVM, however, they are mentioned to lack a deeper semantic meaning. The report points out that transformer-based models such as BERT address these limitations by examining context better, resulting in a higher performance with classification.

The objective of the project is to construct the system to classify news items according to such metrics as accuracy, precision, recall, and F1-score. The desired result is a stable model that will be able to help researchers, journalists, and platforms decrease misinformation and enhance trust in online information.

Inference: Model Transformer-based models such as BERT are much more effective as they use contextual semantics over more traditional feature-based methods [8].

VII. Results

As the experimental assessment shows, the various models perform differently on the fake news dataset, representing a more realistic and challenging classification situation. Conventional machine learning algorithms including Logistic Regression, Naive Bayes and Support Vector machine are very effective in providing a good baseline performance especially in accuracy. But, their accuracy and the ability to memorize are not very high, which indicates the inability to capture complicated linguistic patterns.

Other models such as Ensemble models are better, especially the random Forest, gradient boosting and the XGBoost which demonstrate better performance particularly the recall. This implies that they are effective in detecting the cases of fake news, which is important in reducing false negatives. Of these, Gradient Boosting earns one of the highest recall scores, which implies that it may also be more sensitive to misleading content.

BERT-based model is superior to all classic and ensemble approaches as it is able to capture contextual and semantic connections in the text. In contrast to TF-IDF-based methods, BERT utilizes a bidirection of attention; this allows BERT to comprehend the context of words, too. This leads to both better F1-score and improved precision and the score of recall [10].

The confusion matrix analysis also substantiates that the system produces high true positive rates and reasonable amount of false positives. Although there are still some cases of misclassifications especially in cases of ambiguity, the final results are very high and reflectable in the measures used to evaluate.

Cross-validation Competence: Results of cross-validation show that the model consistently behaves, and there is limited variance in the behaviors of the model using different data splits. This is a confirmation that the models do not over fit and can be extended to unseen data.

Inference: The findings indicate that although classic models offer good baselines, model types that use transformers such as BERT significantly improve the performance because they can detect the deeper contextual information, thus can be more applicable in real-life fake news detection [9].

VIII. Conclusion

By combining machine learning and advanced natural language processing methods, this project is able to build a successful fake news detection system. The system reads text and feeds it to a highly organized pipeline which includes preprocessing, feature extraction, model training and evaluation.

The results emphasize that the dataset is more realistic and complicated, and there are overlapping tendencies in the fake and real news. Due to this, there is no ideal model performance attained by a single traditional model, underscoring the need to select and evaluate models.

The introduction of BERT model will be a great enhancement in the system capacity. Contextual embeddings make BERT more effective in decoding the semantic meaning and identifying fine linguistic words that are related to misinformation. This results in enhanced classification performance and better managing ambiguous content.

Despite the impressive performance of such a system, there are still some difficulties, including how to distinguish the most similar patterns of text and react to the constantly changing tendencies toward misinformation. Such limitations imply that there would be a necessity to constantly update the models and expand the datasets.

On the whole, the proposed system offers an effective and scaled-out solution to the detection of fake news. Such system can be expanded to achieve a set of more powerful capabilities, such as multi-modal integration (text + metadata), real-time deployment, and bigger and more varied datasets to turn it into a comprehensive, misinformation-fighting tool on digital platforms.

Conclusion: The BERT isn't just a feature-based classifier anymore: when incorporated into the system, the BERT transforms it into a context-sensitive intelligent model, enhancing its applicability and reliability in the real world to a larger extent.

REFERENCES

[1] V. Nair, J. Pareek, and S. Bhatt, "A knowledge-based deep learning approach for automatic fake news detection using BERT on Twitter," Procedia Computer Science, vol. 235, pp. 1870–1882, 2024. Available: https://www.sciencedirect.com/science/article/pii/S1877050924008548

[2] M. A. Al Ghamdi, M. S. Bhatti, A. Saeed, Z. Gillani, and S. H. Almotiri, "A fusion of BERT, machine learning and manual approach for fake news detection," *Multimedia Tools and Applications*, vol. 83, no. 10, pp. 30095–30112, 2024. Available: https://link.springer.com/article/10.1007/s11042-023-16669-z

[3] P. Dhiman, A. Kaur, D. Gupta, S. Juneja, A. Nauman, and G. Muhammad, "GBERT: A hybrid deep learning model based on GPT-BERT for fake news detection," *Heliyon*, vol. 10, no. 16, 2024. Available: https://www.cell.com/heliyon/fulltext/S2405-8440(24)11896-8

[4] A. Ramzan, R. H. Ali, N. Ali, and A. Khan, "Enhancing fake news detection using BERT: A comparative analysis of logistic regression, RFC, LSTM and BERT," in *Proceedings of the 2024 International Conference on IT and Industrial Technologies (ICIT)*, pp. 1–6, IEEE, 2024. Available: https://ieeexplore.ieee.org/abstract/document/10859673

[5] S. Mahadevan SR and S. Ahmad, "BERT based blended approach for fake news detection," *Journal of Big Data and Artificial Intelligence*, vol. 2, no. 1, 2024. Available: https://jbdtp.org/index.php/JBDAI/article/view/27

[6] M. Farokhian, V. Rafe, and H. Veisi, "Fake news detection using dual BERT deep neural networks," *Multimedia Tools and Applications*, vol. 83, no. 15, pp. 43831–43848, 2024. Available: https://link.springer.com/article/10.1007/s11042-023-17115-w

[7] M. Farokhian, V. Rafe, and H. Veisi, "Fake news detection using parallel BERT deep neural networks," *arXiv preprint arXiv:2204.04793*, 2022. [Online]. Available: https://arxiv.org/pdf/2204.04793

[8] W. Shishah, "Fake news detection using BERT model with joint learning," *Arabian Journal for Science and Engineering*, vol. 46, no. 9, pp. 9115–9127, 2021. [Online]. Available: https://link.springer.com/article/10.1007/s13369-021-05780-8

[9] M. Szczepański, M. Pawlicki, R. Kozik, and M. Choraś, "New explainability method for BERT-based model in fake news detection," *Scientific Reports*, vol. 11, no. 1, p. 23705, 2021. [Online]. Available: https://www.nature.com/articles/s41598-021-03100-6

[10] J. Ding, Y. Hu, and H. Chang, "BERT-based mental model, a better fake news detector," in *Proceedings of the 2020 6th International Conference on Computing and Artificial Intelligence*, 2020, pp. 396–400. [Online]. Available: https://dl.acm.org/doi/epdf/10.1145/3404555.3404607